

Unit IV

AUTOREGRESSIVE INTEGRATED MOVING AVERAGE (ARIMA) MODELS

Course Outcome:

Explain the exponential smoothing methods
and ARIMA models

What is AUTOREGRESSIVE INTEGRATED MOVING AVERAGE (ARIMA) MODELS

- **ARIMA** (AutoRegressive Integrated Moving Average) models are statistical time series models used to forecast future values of a time series. They are a combination of three components:
 - 1. Autoregressive (AR):** This component assumes that the current value of the series depends on its own past values.
 - 2. Integrated (I):** This component handles non-stationarity in the data by differencing it until it becomes stationary. Differencing involves subtracting the previous value from the current value.
 - 3. Moving Average (MA):** This component assumes that the current value of the series depends on the errors (residuals) from past time periods.

- The ARIMA model is typically denoted as **ARIMA(p, d, q)**, where:
p is the number of lag observations (autoregressive part).
d is the number of times that the raw observations are differenced (integrated part).
q is the size of the moving average window (moving average part).

- **The General Form of an ARIMA Model:**

AR(p): Autoregressive of order p

I(d): Integrated of order d

MA(q): Moving Average of order q

- So, an ARIMA(p,d,q) model means:
- The current value depends on the past p values.
- The data needs to be differenced d times to become stationary.
- The current value depends on the past q error terms.

Why Use ARIMA Models?

- **Forecasting:** They are effective in predicting future values of time series data.
- **Understanding Relationships:** They can help identify patterns and relationships within the data.
- **Modeling Trends and Seasonality:** ARIMA models can capture trends and seasonal components in the data.

Key Considerations:

- **Stationarity:** The data should be stationary (have constant mean and variance) for ARIMA models to be effective.
- **Model Selection:** Choosing the appropriate p , d , and q values is crucial. Techniques like the AIC (Akaike Information Criterion) or BIC (Bayesian Information Criterion) can help in model selection.
- **Model Validation:** Once a model is selected, it's important to validate its performance using techniques like cross-validation.

Applications:

- **Economics:** Forecasting stock prices, GDP, inflation
- **Finance:** Predicting asset returns, risk
- **Marketing:** Sales forecasting, demand analysis
- **Meteorology:** Weather forecasting
- **Environmental Science:** Predicting pollution levels, climate patterns

Linear Models for Stationary Time Series

- **Stationary Time Series**
- A stationary time series is one whose statistical properties (mean, variance, autocorrelation) remain constant over time. In other words, the distribution of the series remains unchanged as time progresses.
- **Key Characteristics of Stationary Time Series:**
- **Constant Mean:** The average value of the series remains the same over time.
- **Constant Variance:** The spread or dispersion of the series remains constant.
- **Constant Autocorrelation:** The relationship between observations at different time points remains consistent.

- **Types of Stationarity:**

- 1. **Strict Stationarity:** The joint probability distribution of any set of observations remains unchanged when shifted in time.

- 2. **Weak Stationarity (Second-Order Stationarity):** The mean, variance, and autocovariance functions remain constant over time.

- **Why is Stationarity Important?**

- **Many statistical methods** assume stationarity for their validity.
- **Stationary time series** are easier to model and forecast.
- **Non-stationary time series** often require transformations or differencing to make them stationary.

- **Examples of Stationary and Non-Stationary Time Series:**
- **Stationary:** Daily temperature readings, monthly stock returns (after accounting for inflation)
- **Non-Stationary:** GDP growth rates, stock prices (often exhibit trends or random walks)

Linear Models for Stationary Time Series

- **Linear models** are a fundamental class of models used to describe the relationship between a dependent variable and one or more independent variables. In the context of stationary time series, these models are used to predict future values based on past observations.

Common Linear Models for Stationary Time Series

1. Autoregressive (AR) Model:

1. Definition: An AR model assumes that the current value of a time series depends linearly on its own past values.

2. Equation:

$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + \varepsilon_t$$

Parameters:

- c : Intercept
- $\phi_1, \phi_2, \dots, \phi_p$: Autoregressive coefficients
- ε_t : White noise error term

- Using the **ar** Function from Base R.
- The **ar function** in base R can be used to fit an autoregressive model to a time series.
- **Syntax:**
- **ar(x, order.max = NULL, aic = TRUE,
method = c("yule-walker", "burg", "ols", "mle"))**
 - x: A numeric vector or time series.
 - order.max: The maximum number of lags to consider.
 - aic: Whether to select the model based on AIC.
 - method: The method to use for fitting (e.g., "yule-walker", "burg", "ols", "mle").

Example

```
# Load necessary library  
library(stats)
```

```
# Generate some example data  
set.seed(123)  
data <- arima.sim(n = 100, model = list(ar = c(0.7)))
```

```
# Fit an AR model  
ar_model <- ar(data, order.max = 10)
```

```
# Print the AR model details  
print(ar_model)
```

Using the Arima Function from the forecast Package

- The Arima function from the forecast package can fit AR, MA, and ARMA models and allows for more flexibility.
- Syntax:

`Arima(y, order = c(p, d, q), seasonal = list(order = c(P, D, Q), period = s))`

- `y`: The time series data.
- `order`: The non-seasonal part of the model (p, d, q).
- `seasonal`: A list specifying the seasonal part of the model.

```
# Install and load the forecast package  
install.packages("forecast")  
library(forecast)
```

```
# Generate some example data  
set.seed(123)  
data <- arima.sim(n = 100, model = list(ar = c(0.7)))
```

```
# Fit an AR(1) model  
ar_model <- Arima(data, order = c(1, 0, 0))
```

```
# Print the AR model details  
print(ar_model)
```

Moving Average (MA) Model

- **Definition:** The MA model expresses the current value of the series as a linear combination of past white noise error terms.
- Equation:

$$y_t = c + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \dots + \theta_q \varepsilon_{t-q} + \varepsilon_t$$

Parameters:

- c : Intercept
- $\theta_1, \theta_2, \dots, \theta_q$: Moving average coefficients
- ε_t : White noise error term

- Using the `arima` Function from Base R.
- The `arima` function in base R can be used to fit a Moving Average (MA) model as part of an ARMA or ARIMA model.
- To fit just an MA model, you set the autoregressive (p) and differencing (d) components to zero.
- **Syntax:**

`arima(y, order = c(p, d, q), seasonal = list(order = c(P, D, Q), period = s))`

- `y`: The time series data.
- `order`: A vector specifying the non-seasonal part of the model (p, d, q).
- `p`: The number of autoregressive terms (set to 0 for MA models).
- `d`: The number of differences (set to 0 for MA models).
- `q`: The number of moving average terms.
- `seasonal`: Optional list specifying seasonal components.

Example

```
# Generate some example data
set.seed(123)
data <- arima.sim(n = 100, model = list(ma = c(0.5)))

# Fit an MA(1) model
ma_model <- arima(data, order = c(0, 0, 1))

# Print the MA model details
print(ma_model)
```

- Using the Arima Function from the forecast Package
- The Arima function from the forecast package is flexible and can fit MA models as part of ARMA or ARIMA models.

```
# Install and load the forecast package
```

```
install.packages("forecast")
```

```
library(forecast)
```

```
# Generate some example data
```

```
set.seed(123)
```

```
data <- arima.sim(n = 100, model = list(ma = c(0.5)))
```

```
# Fit an MA(1) model
```

```
ma_model <- Arima(data, order = c(0, 0, 1))
```

```
# Print the MA model details
```

```
print(ma_model)
```

Autoregressive Moving Average (ARMA) Model:

- **Definition:** An ARMA model combines the AR and MA models, assuming that the current value of a time series depends on both its past values and past error terms.

Equation:

$$y_t = c + \varphi_1 y_{(t-1)} + \varphi_2 y_{(t-2)} + \dots + \varphi_p y_{(t-p)} + \theta_1 \varepsilon_{(t-1)} + \theta_2$$

Parameters:

- c : Intercept
- $\varphi_1, \varphi_2, \dots, \varphi_p$: Autoregressive coefficients
- $\theta_1, \theta_2, \dots, \theta_q$: Moving average coefficients
- ε_t : White noise error term

Autoregressive Moving Average (ARMA) Model:

- Using the `arima` Function from Base R
- The `arima` function from base R can fit ARMA models by specifying the order of autoregressive (p) and moving average (q) components, with differencing (d) set to zero.
- **Syntax:**
`arima(y, order = c(p, d, q), seasonal = list(order = c(P, D, Q), period = s))`
 - `y`: The time series data.
 - `order`: A vector specifying the non-seasonal part of the model (p, d, q).
 - `p`: The number of autoregressive terms.
 - `d`: The number of differences (set to 0 for ARMA models).
 - `q`: The number of moving average terms.
 - `seasonal`: Optional list specifying seasonal components.

Example

```
# Generate some example data
set.seed(123)
data <- arima.sim(n = 100, model = list(ar = c(0.7), ma = c(0.5)))

# Fit an ARMA(1,1) model
arma_model <- arima(data, order = c(1, 0, 1))

# Print the ARMA model details
print(arma_model)
```

Model Selection and Estimation

- **Stationarity:** Ensure the time series is stationary before applying these models.
- **Model Identification:** Use tools like the autocorrelation function (ACF) and partial autocorrelation function (PACF) to identify potential AR and MA orders.
- **Model Estimation:** Employ methods like least squares or maximum likelihood to estimate the model parameters.
- **Model Validation:** Assess the model's performance using techniques like cross-validation or hypothesis testing.

Finite Order Moving Average Processes

- **Finite Order Moving Average (MA) Processes** are a class of time series models that assume the current value of a series depends linearly on the current and past error terms. The "finite order" part indicates that the model considers only a finite number of past error terms.

- **General Form of an MA(q) Process:**

$$y_t = \mu + \varepsilon_t + \theta_1\varepsilon_{(t-1)} + \theta_2\varepsilon_{(t-2)} + \dots + \theta_q\varepsilon_{(t-q)}$$

- y_t : The value of the time series at time t
- μ : The mean of the series
- ε_t : White noise error term at time t
- $\theta_1, \theta_2, \dots, \theta_q$: Moving average coefficients
- q : The order of the MA process

Finite Order Moving Average Processes

- **Key Characteristics of MA Processes:**
- **Stationarity:** MA processes are inherently stationary, meaning their statistical properties (mean, variance, autocorrelation) remain constant over time.
- **Autocorrelation:** The autocorrelation function (ACF) of an $MA(q)$ process cuts off after lag q . This means that the correlation between observations more than q periods apart is zero.
- **Invertibility:** An MA process is invertible if it can be expressed as an equivalent AR process. Invertibility ensures that the model can be used for forecasting.

Finite Order Moving Average Processes

- **Applications of MA Processes:**
- **Forecasting:** MA models are useful for short-term forecasting, especially when the series exhibits a relatively simple pattern.
- **Noise Reduction:** MA models can be used to remove noise from time series data.
- **Signal Processing:** MA processes are used in signal processing applications, such as filtering and smoothing.

Finite Order Moving Average Processes

- **Example**
- **1. Simulating a Finite Order MA Process**
- We'll simulate a Moving Average (MA) process of order 2, denoted $MA(2)$. In this process, the current value depends on the current and previous two white noise error terms.
- **Steps:**
 1. Simulate white noise error terms.
 2. Generate the $MA(2)$ time series using these error terms.

```
# Set seed for reproducibility
set.seed(123)
# Parameters
n <- 100 # Number of observations
theta1 <- 0.5 # MA(1) coefficient
theta2 <- -0.3 # MA(2) coefficient

# Simulate white noise error terms
epsilon <- rnorm(n + 2) # Extra terms to handle initial conditions

# Generate MA(2) process
ma2_data <- filter(epsilon, filter = c(1, -theta1, -theta2), method = "convolution", sides = 1)

# The MA(2) process starts from the third value because of filter initialization
ma2_data <- ma2_data[-(1:2)]

# Plot the simulated MA(2) data
plot(ma2_data, type = "l", main = "Simulated MA(2) Process", ylab = "Value", xlab = "Time")
```

Finite Order Moving Average Processes

- 2. Fitting an MA(2) Model To fit an MA(2) model to the simulated data, we use the `arima` function from base R, setting the autoregressive (p) and differencing (d) components to zero.

```
# Fit an MA(2) model
```

```
fit <- arima(ma2_data, order = c(0, 0, 2))
```

```
# Print the MA model details
```

```
print(fit)
```

- **3. Analyzing Residuals**

- After fitting the model, it's crucial to analyze the residuals to ensure they resemble white noise. We'll plot the residuals and check their autocorrelation.

- **R Code:**

```
# Plot the residuals
```

```
residuals <- residuals(fit)
```

```
plot(residuals, type = "l", main = "Residuals of the MA(2) Model", ylab =  
"Residuals", xlab = "Time")
```

```
# Plot the ACF of the residuals
```

```
acf(residuals, main = "ACF of Residuals")
```

```
# Plot the PACF of the residuals
```

```
pacf(residuals, main = "PACF of Residuals")
```

Finite Order Moving Average Processes

Summary

Here's what each part of the example does:

- **Simulating MA(2):** Uses filter to generate an MA(2) process from white noise.
- **Fitting MA(2) Model:** Uses arima to estimate an MA(2) model and prints the details.
- **Analyzing Residuals:** Plots residuals and their ACF/PACF to ensure the model is a good fit.

Finite Order Autoregressive Processes

Finite Order Autoregressive (AR) Processes are a class of time series models that assume the current value of a series depends linearly on its own past values. The "finite order" part indicates that the model considers only a finite number of past values.

General Form of an AR(p) Process:

$$y_t = \mu + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + \varepsilon_t$$

- y_t : The value of the time series at time t
- μ : The mean of the series
- $\phi_1, \phi_2, \dots, \phi_p$: Autoregressive coefficients
- ε_t : White noise error term
- p : The order of the AR process

Finite Order Autoregressive Processes

Key Characteristics of AR Processes:

- **Stationarity:** AR processes are inherently stationary under certain conditions (e.g., when the autoregressive coefficients are within the unit circle).
- **Autocorrelation:** The autocorrelation function (ACF) of an AR(p) process decays exponentially or geometrically. This means that the correlation between observations at distant time points gradually decreases.
- **Partial Autocorrelation Function (PACF):** The PACF of an AR(p) process cuts off after lag p. This means that the correlation between y_t and $y_{(t-k)}$ for $k > p$ is zero when controlling for the intervening values.

- **Applications of AR Processes:**
- **Forecasting:** AR models are useful for forecasting future values of a time series, especially when the series exhibits a strong dependence on its past values.
- **Modeling Trends and Seasonality:** AR models can be used to capture trends and seasonal patterns in time series data.
- **Control Systems:** AR models are used in control systems to predict future system states and adjust inputs accordingly.

Example

- **1. Simulating a Finite Order AR Process**

- We'll simulate an Autoregressive (AR) process of order 2, denoted AR(2). In this process, the value at time t is influenced by the values at times $t-1$ and $t-2$.

- **Steps:**

1. Generate the AR(2) time series using predefined AR coefficients.
2. Plot the simulated data.

Example

```
# Set seed for reproducibility
set.seed(123)
# Parameters
n <- 100 # Number of observations
phi1 <- 0.6 # AR(1) coefficient
phi2 <- -0.2 # AR(2) coefficient
# Generate AR(2) process
# Coefficients for AR(2) process: phi1, phi2
# Use arima.sim to generate the process
ar2_data <- arima.sim(n = n, model = list(ar = c(phi1, phi2)))

# Plot the simulated AR(2) data
plot(ar2_data, type = "l", main = "Simulated AR(2) Process", ylab = "Value", xlab = "Time")
```

2. Fitting an AR(2) Model

To fit an AR(2) model to the simulated data, use the `arima` function from base R, specifying the order of the AR component.

```
# Fit an AR(2) model
ar2_fit <- arima(ar2_data, order = c(2, 0, 0))

# Print the AR model details
print(ar2_fit)
```

- **3. Analyzing Residuals**

- After fitting the model, analyze the residuals to ensure they resemble white noise. This involves plotting the residuals and checking their autocorrelation.

```
# Plot the residuals
```

```
residuals <- residuals(ar2_fit)
```

```
plot(residuals, type = "l", main = "Residuals of the AR(2) Model", ylab =  
"Residuals", xlab = "Time")
```

```
# Plot the ACF of the residuals
```

```
acf(residuals, main = "ACF of Residuals")
```

```
# Plot the PACF of the residuals
```

```
pacf(residuals, main = "PACF of Residuals")
```

Summary

Here's a breakdown of the steps:

- **Simulating AR(2) Process:** Uses `arima.sim` to generate AR(2) data based on specified coefficients.
- **Fitting AR(2) Model:** Uses `arima` to fit an AR(2) model to the simulated data.
- **Analyzing Residuals:** Plots residuals and their ACF/PACF to ensure that residuals are white noise.

Mixed Autoregressive–Moving Average Processes

Mixed Autoregressive–Moving Average (ARMA) Processes are a combination of Autoregressive (AR) and Moving Average (MA) processes. They are used to model time series data that exhibit both autoregressive and moving average characteristics.

Mixed Autoregressive–Moving Average Processes

General Form of an ARMA(p, q) Process:

$$y_t = \mu + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \dots + \theta_q \varepsilon_{t-q} + \varepsilon_t$$

- y_t : The value of the time series at time t
- μ : The mean of the series
- $\phi_1, \phi_2, \dots, \phi_p$: Autoregressive coefficients
- $\theta_1, \theta_2, \dots, \theta_q$: Moving average coefficients
- ε_t : White noise error term
- p : The order of the AR component
- q : The order of the MA component

Mixed Autoregressive–Moving Average Processes

- **Key Characteristics of ARMA Processes:**
- **Stationarity:** ARMA processes are stationary under certain conditions (e.g., when the AR and MA coefficients are within the unit circle).
- **Autocorrelation(ACF) and Partial Autocorrelation(PACF):** The ACF and PACF of an ARMA process exhibit a combination of exponential decay (from the AR component) and cutting off (from the MA component).
- **Invertibility:** An ARMA process is invertible if it can be expressed as an equivalent AR process. Invertibility ensures that the model can be used for forecasting.

Mixed Autoregressive–Moving Average Processes

- **Applications of ARMA Processes:**
- **Forecasting:** ARMA models are widely used for forecasting future values of time series data.
- **Modeling Trends and Seasonality:** ARMA models can capture both trends and seasonal patterns.
- **Noise Reduction:** ARMA models can be used to remove noise from time series data.

Example

- **1. Simulating an ARMA(2,1) Process**

- An ARMA(2,1) process has:

- AR terms: $p=2$ (two past values)

- MA terms: $q=1$ (one past error term)

- **Steps:**

1. Generate the ARMA(2,1) time series using predefined AR and MA coefficients.

2. Plot the simulated data.

```
# Set seed for reproducibility  
set.seed(123)
```

```
# Parameters  
n <- 100 # Number of observations  
phi1 <- 0.6 # AR(1) coefficient  
phi2 <- -0.2 # AR(2) coefficient  
theta1 <- 0.5 # MA(1) coefficient
```

```
# Generate ARMA(2,1) process  
# arima.sim function can be used to simulate ARMA processes  
arma21_data <- arima.sim(n = n, model = list(ar = c(phi1, phi2), ma = c(theta1)))
```

```
# Plot the simulated ARMA(2,1) data  
plot(arma21_data, type = "l", main = "Simulated ARMA(2,1) Process", ylab = "Value",  
xlab = "Time")
```

2. Fitting an ARMA(2,1) Model

To fit an ARMA(2,1) model to the simulated data, use the `arima` function from base R. Specify the order of AR and MA components accordingly.

R Code:

```
# Fit an ARMA(2,1) model  
arma21_fit <- arima(arma21_data, order = c(2, 0, 1))  
  
# Print the ARMA model details  
print(arma21_fit)
```

- **3. Analyzing Residuals**

- After fitting the model, it's important to analyze the residuals to ensure they resemble white noise. This involves plotting the residuals and checking their autocorrelation.

```
# Extract residuals from the fitted model  
residuals_arma21 <- residuals(arma21_fit)
```

```
# Plot the residuals  
plot(residuals_arma21, type = "l", main = "Residuals of the ARMA(2,1) Model",  
ylab = "Residuals", xlab = "Time")
```

```
# Plot the ACF of the residuals  
acf(residuals_arma21, main = "ACF of Residuals")
```

```
# Plot the PACF of the residuals  
pacf(residuals_arma21, main = "PACF of Residuals")
```

Mixed Autoregressive–Moving Average Processes

Summary

Here's a breakdown of the steps:

- **Simulating ARMA(2,1) Process:** Uses `arima.sim` to generate ARMA(2,1) data based on specified AR and MA coefficients.
- **Fitting ARMA(2,1) Model:** Uses `arima` to fit an ARMA(2,1) model to the simulated data.
- **Analyzing Residuals:** Plots residuals and their ACF/PACF to ensure that residuals are white noise.

Mixed Autoregressive Nonstationary Processes

- Mixed Autoregressive Nonstationary Processes (MANPs) are a class of time series models that combine autoregressive (AR) components with nonstationary characteristics.
- These models are particularly useful for analyzing time series data where the statistical properties, such as mean and variance, change over time.

- **Key Concepts:**
- **Autoregressive (AR) Processes:** These are models where the current value of the series is based on its previous values. For example, an AR(1) process can be written as:
 - $X_t = \phi X_{t-1} + \epsilon_t$
 - where $(\phi \ \backslash \text{phi})$ is the autoregressive coefficient and $(\epsilon \backslash \text{epsilon}_t)$ is white noise.
- **Nonstationarity:** This refers to time series whose statistical properties change over time. Nonstationary processes can exhibit trends, seasonal effects, or other time-varying structures.
- **Mixed Processes:** MANPs combine the autoregressive structure with nonstationary components, allowing for more flexible modeling of complex time series data.

- **Applications:**
- **Economics:** Modeling economic indicators that evolve over time.
- **Finance:** Analyzing stock prices or interest rates that exhibit nonstationary behavior.
- **Environmental Science:** Studying climate data with changing patterns over time.

- **Summary**
- Working with mixed autoregressive nonstationary processes involves:
- **Exploring and preparing your data**
- **Testing for and addressing nonstationarity**
- **Fitting appropriate autoregressive models**
- **Handling complex dynamics with advanced models if necessary**
- **Diagnosing and validating your model**
- **Making forecasts and evaluating performance**

- **1. Install Necessary Packages**

```
install.packages(c("forecast", "tseries", "NTS"))  
library(forecast)  
library(tseries)  
library(NTS)
```

- **2. Data Preparation**

```
data <- ts(your_data, start = c(start_year, start_period), frequency = frequency)
```

- **3. Stationarity Check**

- Check if your time series is stationary. If not, you may need to difference the data.

```
adf.test(data) # Augmented Dickey-Fuller test  
data_diff <- diff(data)
```

4. Model Fitting

Fit an autoregressive model to your data. You can use the `arima` function for ARIMA models or the `sarima` function from the `astsa` package for seasonal ARIMA models.

```
model <- arima(data_diff, order = c(p, d, q))
```

5. Nonlinear Models

For more complex models, the `NTS` package provides tools for nonlinear time series analysis, including threshold autoregressive (TAR) models and state-space models.

```
# Example of fitting a TAR model  
tar_model <- tar(data_diff, p = 1, d = 1, q = 1)
```

6. Model Evaluation

Evaluate the model's performance using various diagnostic tools.

```
tsdiag(model)
```

Example

- `# Load necessary libraries`
- `library(ggplot2)`
- `library(forecast)`
- `library(tseries)`
- `# Set random seed for reproducibility`
- `set.seed(123)`
- `# Generate synthetic data`
- `n <- 200`
- `time_series <- 0.5 * (1:n) + arima.sim(model = list(ar = 0.7), n = n) +
rnorm(n, sd = 2) ts_data <- ts(time_series, frequency = 12)`

- `# Plot the synthetic time series`
- `autoplot(ts_data) + ggtitle("Synthetic Time Series with Trend and AR Component")`
- `# Perform Augmented Dickey-Fuller test`
- `adf_result <- adf.test(ts_data, alternative = "stationary")`
`print(adf_result)`
- `# Difference the time series`
- `diff_ts_data <- diff(ts_data)`
- `# Plot differenced time series`
- `autoplot(diff_ts_data) + ggtitle("Differenced Time Series")`

- # Check for stationarity again
- `adf_result_diff <- adf.test(diff_ts_data, alternative = "stationary")`
`print(adf_result_diff)`
- # Fit ARIMA model to the original data
- `fit <- auto.arima(ts_data)` `summary(fit)`
- # Forecast using the ARIMA model
- `forecasted <- forecast(fit, h = 12)`
- `autoplot(forecasted) + ggtitle("ARIMA Model Forecast")`
- # Residual diagnostics `checkresiduals(fit)`

Seasonal ARIMA (SARIMA)

- Seasonal ARIMA (SARIMA) models are an extension of the ARIMA (AutoRegressive Integrated Moving Average) models, specifically designed to handle seasonality in time series data. They are useful for forecasting when data exhibit seasonal patterns, such as monthly sales figures or quarterly economic indicators.

Components of Seasonal ARIMA (SARIMA)

- **Non-Seasonal Part:**

- **AR (p):** Autoregressive terms.
- **I (d):** Differencing to make the series stationary.
- **MA (q):** Moving average terms.

- **Seasonal Part:**

- **Seasonal AR (P):** Seasonal autoregressive terms.
- **Seasonal I (D):** Seasonal differencing.
- **Seasonal MA (Q):** Seasonal moving average terms.
- **m:** Number of periods per season (e.g., 12 for monthly data with yearly seasonality).

- **Model Notation**
- SARIMA models are denoted as:
- $\text{ARIMA}(p,d,q)(P,D,Q)_m$

- **Steps to Build a SARIMA Model**
- **Identify Seasonality:** Determine the seasonal period (m).
- **Differencing:** Apply seasonal differencing to remove seasonality and non-seasonal differencing to remove trends.
- **Model Selection:** Use ACF and PACF plots to identify potential AR and MA terms.
- **Parameter Estimation:** Estimate the parameters using statistical software.
- **Model Diagnostics:** Check residuals to ensure they resemble white noise.

Applications

- **1. Economics and Finance**
- **Sales Forecasting:** Predicting future sales based on past seasonal patterns, such as holiday sales spikes.
- **Stock Market Analysis:** Analyzing and forecasting stock prices that exhibit seasonal trends.
- **Economic Indicators:** Forecasting economic indicators like GDP, inflation rates, and unemployment rates that have seasonal components.
- **2. Environmental Science**
- **Weather Forecasting:** Predicting seasonal weather patterns, such as temperature and precipitation.
- **Air Quality Monitoring:** Forecasting pollution levels that vary with seasons.
- **Hydrology:** Predicting river flows and water levels that are influenced by seasonal rainfall and snowmelt.
- **3. Retail and Inventory Management**
- **Demand Forecasting:** Anticipating product demand to manage inventory levels, especially for seasonal products.
- **Supply Chain Optimization:** Planning and optimizing supply chain operations based on seasonal demand patterns.

Applications

- **4. Healthcare**
 - **Disease Outbreak Prediction:** Forecasting the spread of seasonal diseases like influenza.
 - **Hospital Resource Management:** Predicting patient admissions and resource needs based on seasonal trends.
- **5. Tourism and Hospitality**
 - **Visitor Forecasting:** Predicting tourist arrivals and hotel bookings that vary with seasons.
 - **Event Planning:** Planning events and promotions based on expected seasonal demand.
- **6. Agriculture**
 - **Crop Yield Prediction:** Forecasting crop yields that are influenced by seasonal weather patterns.
 - **Market Prices:** Predicting seasonal fluctuations in the prices of agricultural products.
- **7. Energy Sector**
 - **Electricity Demand:** Forecasting electricity consumption that varies with seasons, such as higher demand in summer for cooling and in winter for heating.
 - **Renewable Energy Production:** Predicting the output of renewable energy sources like solar and wind, which are affected by seasonal weather patterns.

Time Series Model Building Process

- **1. Data Collection and Preprocessing**
 - **Collect Data:** Gather historical time series data relevant to the problem.
 - **Clean Data:** Handle missing values, outliers, and any inconsistencies in the data.
 - **Transform Data:** Apply necessary transformations (e.g., log transformation) to stabilize variance.
- **2. Exploratory Data Analysis (EDA)**
 - **Plot the Data:** Visualize the time series to identify trends, seasonality, and any irregular patterns.
 - **Decompose the Series:** Decompose the time series into trend, seasonal, and residual components.
 - **Summary Statistics:** Calculate summary statistics to understand the data distribution.
- **3. Stationarity Check**
 - **Visual Inspection:** Check if the mean and variance are constant over time.
 - **Statistical Tests:** Use tests like the Augmented Dickey-Fuller (ADF) test to check for stationarity.
 - **Differencing:** Apply differencing to make the series stationary if needed.
- **4. Model Identification**
 - **ACF and PACF Plots:** Analyze Autocorrelation Function (ACF) and Partial Autocorrelation Function (PACF) plots to identify potential AR and MA terms.
 - **Model Selection:** Choose the appropriate model (e.g., ARIMA, SARIMA) based on the data characteristics.

Time Series Model Building Process

- **Parameter Estimation**
 - **Fit the Model:** Use statistical software to estimate the model parameters.
 - **Optimize Parameters:** Fine-tune the parameters to improve model performance.
- **6. Model Diagnostics**
 - **Residual Analysis:** Check the residuals to ensure they resemble white noise (i.e., no autocorrelation).
 - **Diagnostic Plots:** Use plots like Q-Q plots and residual plots to assess model fit.
 - **Statistical Tests:** Perform tests like the Ljung-Box test to check for autocorrelation in residuals.
- **7. Model Validation**
 - **Train-Test Split:** Split the data into training and testing sets.
 - **Cross-Validation:** Use techniques like rolling cross-validation to validate the model.
 - **Performance Metrics:** Evaluate the model using metrics like Mean Absolute Error (MAE), Mean Squared Error (MSE), and Root Mean Squared Error (RMSE).
- **8. Forecasting**
 - **Generate Forecasts:** Use the model to make future predictions.
 - **Confidence Intervals:** Provide confidence intervals for the forecasts to quantify uncertainty.
 - **Visualize Forecasts:** Plot the forecasts along with the historical data for comparison.

- **9. Model Deployment**

- **Implement the Model:** Deploy the model in a production environment.
- **Monitor Performance:** Continuously monitor the model's performance and update it as needed.

- **10. Model Maintenance**

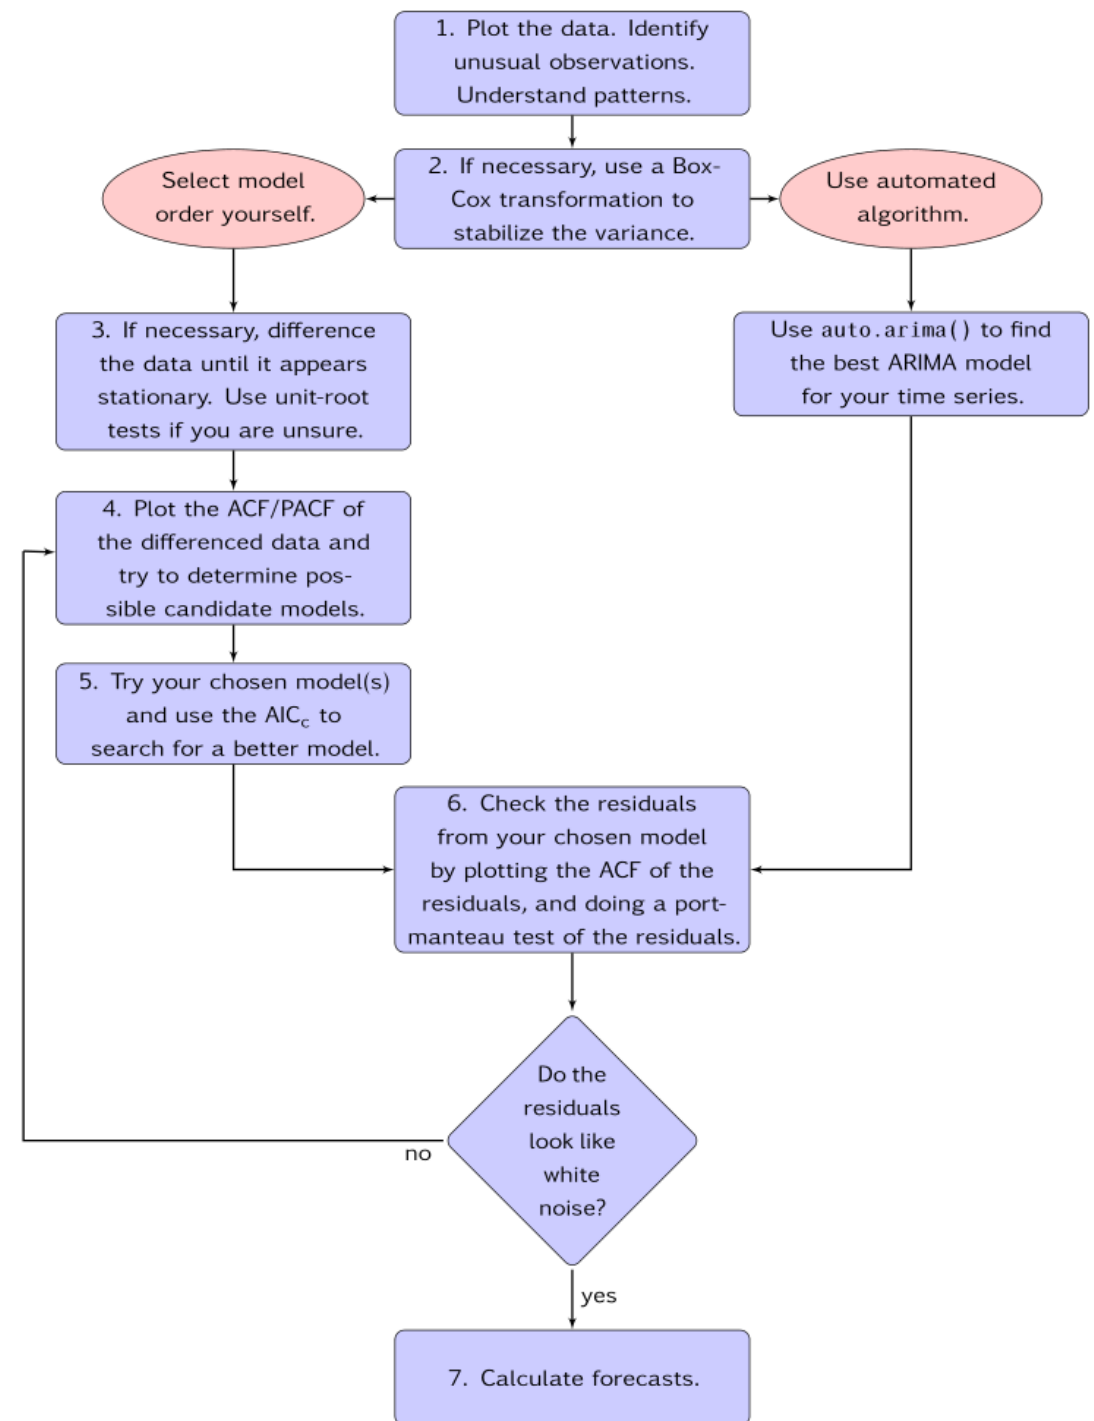
- **Retrain the Model:** Periodically retrain the model with new data to maintain accuracy.
- **Update Parameters:** Adjust model parameters based on new insights and data patterns.

This process ensures a systematic approach to building robust time series models. If you need more details on any specific step or have a particular dataset in mind, feel free to ask!

Modelling procedure

General process for forecasting using an ARIMA model.

1. Plot the data and identify any unusual observations.
2. If necessary, transform the data (using a Box-Cox transformation) to stabilise the variance.
3. If the data are non-stationary, take first differences of the data until the data are stationary.
4. Examine the ACF/PACF: Is an $\text{ARIMA}(p,d,0p,d,0)$ or $\text{ARIMA}(0,d,q0,d,q)$ model appropriate?
5. Try your chosen model(s), and use the AIC_c to search for a better model.
6. Check the residuals from your chosen model by plotting the ACF of the residuals, and doing a portmanteau test of the residuals. If they do not look like white noise, try a modified model.
7. Once the residuals look like white noise, calculate forecasts.



Modeling bio-surveillance data using ARIMA (AutoRegressive Integrated Moving Average)

- Modeling bio-surveillance data using ARIMA (AutoRegressive Integrated Moving Average) is a powerful approach for forecasting and analyzing time series data, particularly in the context of public health and epidemiology.

Key Steps in ARIMA Modeling for Bio-surveillance Data:

1. Data Collection:

- Gather historical bio-surveillance data, such as disease incidence rates, hospital admissions, or other relevant health metrics.

2. Data Preprocessing:

- **Stationarity:** Ensure the time series data is stationary, meaning its statistical properties do not change over time. This often involves differencing the data.
- **Seasonality:** Identify and handle any seasonal patterns in the data, which might require using Seasonal ARIMA (SARIMA).

3. Model Identification:

- Determine the order of the ARIMA model, denoted as (p, d, q) :
 - **p:** Number of lag observations (autoregressive part).
 - **d:** Degree of differencing (to make the series stationary).
 - **q:** Size of the moving average window.

4. Parameter Estimation:

- Use statistical software or programming languages like Python (with libraries such as statsmodels) to estimate the parameters of the ARIMA model.

5. Model Diagnostics:

- Check the residuals of the model to ensure they resemble white noise, indicating a good fit.

6. Forecasting:

- Use the fitted ARIMA model to make forecasts and predict future values of the bio-surveillance data.

Example in Python: Here's a basic example of how you might implement ARIMA modeling in Python:

```
import pandas as pd
from statsmodels.tsa.arima.model import ARIMA
import matplotlib.pyplot as plt
```

```
# Load your bio-surveillance data
data = pd.read_csv('bio_surveillance_data.csv')
data.index = pd.to_datetime(data['date'])
series = data['incidence_rate']
```

```
# Fit the ARIMA model
model = ARIMA(series, order=(p, d, q))
model_fit = model.fit()
```

```
# Summary of the model
print(model_fit.summary())
```

```
# Forecasting
forecast = model_fit.forecast(steps=10)
plt.plot(series, label='Observed')
plt.plot(forecast, label='Forecast', color='red')
plt.legend()
plt.show()
```

Applications:

- Epidemiological Surveillance: Forecasting the spread of infectious diseases¹.
- Public Health Planning: Predicting hospital admissions and resource needs².
- Environmental Health: Monitoring air quality and predicting pollution levels

R Commands

Step 1: Install and Load Required Packages

```
install.packages("forecast") library(forecast)
```

Step 2: Load Your Data

Assuming your data is a time series object:

```
# Example: Load data (replace this with your actual data) data <- ts(your_data_vector,  
frequency = 12) # Monthly data
```

Step 3: Visualize the Data

```
plot(data, main = "Time Series Data", ylab = "Values", xlab = "Time")
```

Step 4: Check for Stationarity

You can use the Augmented Dickey-Fuller test:

```
library(tseries) adf.test(data)
```


Step 5: Differencing if Necessary

If the data is not stationary, apply differencing:

```
diff_data <- diff(data, differences = 1)
# Non-seasonal differencing seasonal_diff_data <- diff(data, lag = 12)
# Seasonal differencing
```

Step 6: Identify Parameters (p, d, q)

You can analyze ACF and PACF plots:

```
acf(data)      # Autocorrelation Function pacf(data) # Partial Autocorrelation Function
```

Step 7: Fit the ARIMA Model

Use the `auto.arima` function for automatic selection of parameters:

```
fit <- auto.arima(data) summary(fit)
```

For a specific SARIMA model, use:

```
fit <- Arima(data, order = c(p, d, q), seasonal = c(P, D, Q))
```

Step 8: Diagnostic Checks

Check the residuals of the fitted model:

```
checkresiduals(fit)
```

Step 9: Forecasting

Make forecasts with the fitted model:

```
forecasts <- forecast(fit, h = 12) # Forecast for the next 12 periods  
plot(forecasts)
```

Step 10: Evaluate the Model

You can evaluate the model using various accuracy metrics:

```
accuracy(forecasts)
```

Example

Load necessary libraries

```
install.packages("forecast")  
library(forecast)  
library(tseries)
```

Load and prepare data

```
data <- ts(your_data_vector, frequency = 12)
```

Visualize the data

```
plot(data)
```

Check for stationarity

```
adf.test(data)
```

Fit the SARIMA model

```
fit <- auto.arima(data)
```

Check model summary

```
summary(fit)
```

Diagnostic checks

```
checkresiduals(fit)
```

Forecasting

```
forecasts <- forecast(fit, h = 12)  
plot(forecasts)
```

Evaluate the forecast

```
accuracy(forecasts)
```